

CS553 Programming Assignment #1

Benchmarking

TEAM MEMBERS:

- **VAISHNAVI GUDUR (A20384900)**
- **SATYA SUVEEN REDDY MEKALA(A20379568)**

DESIGN DOCUMENT

The assignment involves benchmarking different parts of a computer system, from the CPU, GPU, Memory, Disk, and Network.

The basic design and logic implemented for each of the above mentioned 5 components is explained below:

CPU BENCHMARKING:

1. Program Design:

The main goal of this experiment is to obtain the CPU speed in Giga IOPS and Giga FLOPS by performing arithmetic operations using Integers and Double precision floating point numbers respectively, with varying levels of concurrency.

- The code is designed in such a way that there are 3 methods defined including the main function; main function defines the number of threads (1,2,4 and 8) and calls the second function for each of the number of threads.
- The second function includes thread creation, pthread_join() function suspends the execution of calling thread unless the target thread is terminated. Also calculates the time required to perform the operations and processor speed.
- The third function performs the arithmetic operations.
- To calculate IOPS and FLOPS, 30 arithmetic operations involving addition and subtraction are performed for 10^9 times on both integers and double precision floating point numbers.
- The time required to complete the execution is recorded at each concurrency level by recording the start time at the start of the operation and end time after completion of the operation, and calculating the difference thereafter.
- In case of more than 1 thread, the total time of execution of all the threads is recorded.
- The speed in IOPS and GFLOPS is calculated as the product of number of times the operations performed times the number of arithmetic operations performed divided by the execution time elapsed after performing those operations.
- The same program is run with 1 thread, 2 threads, 4 threads and 8 threads.

- The program is designed to support strong scaling; the number of instructions is kept fixed for the entire code. However, as the number of threads increase, the amount of work per thread is decreased such that each thread performs equal share of work.

2. Performance:

The following formula was used to calculate the processor speed in GIOPS and GFLOPS:

$$\text{Speed} = \frac{\text{Number of arithmetic instructions} * \text{Number of times operations performed}}{\text{Execution time} * 10^9}$$

Theoretical Speed:

The theoretical speed of CPU can be calculated as,

$$\text{GHz} * \text{Cores} * \text{IPC}.$$

Chameleon Openstack instance has **2.29 GHz, 2 Cores and 16 IPC for Intel Core Processor (Haswell)**. I obtained this information using ‘lscpu’ command on the instance.

Since the processor info was stated as Intel core(Haswell) of family 6 Series 60 stepping 1, the instruction per cycle was decided as 8 IPC .

$$\begin{aligned} \text{Thus, theoretical peak performance for 2 threads (1 thread/core)} &= 2.29 * 2 * 8 \\ &= \mathbf{36.64 \text{ GFLOPS}} \end{aligned}$$

The results obtained after repeatedly performing the experiment are as follows:

GIOPS:

Sr. No.	1 thread		2 threads		4 threads		8 threads	
	GIOPS	TIME	GIOPS	TIME	GIOPS	TIME	GIOPS	TIME
1.	4.189944	7.16	4.010695	7.48	4.195804	7.15	4.137931	7.25
2.	4.231312	7.09	4.160888	7.21	4.149378	7.23	4.243281	7.07
3.	4.297994	6.98	4.304161	6.97	4.115226	7.29	4.137931	7.25

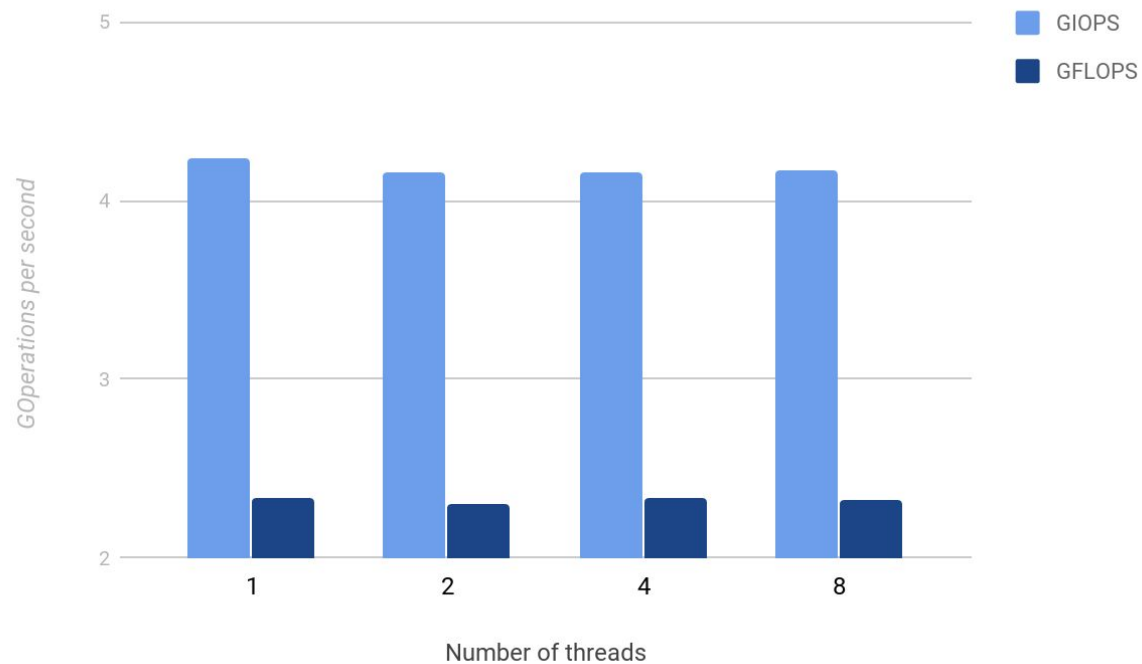
Average	4.23975	7.08	4.158581	7.22	4.153469	7.22	4.173048	7.19
Standard Deviation	0.044513	0.07	0.119818	0.21	0.033023	0.06	0.049662	0.08

GFLOPS:

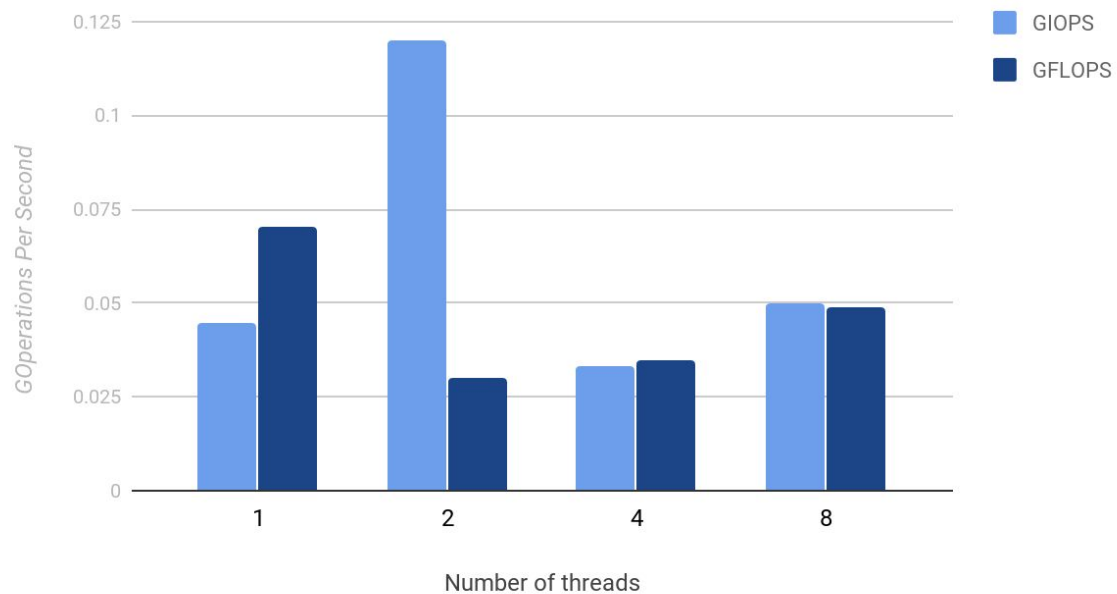
Sr. No.	1 thread		2 threads		4 threads		8 threads	
	GFLOPS	TIME	GFLOPS	TIME	GFLOPS	TIME	GFLOPS	TIME
1.	2.327386	12.89	2.305919	13.01	2.382844	12.59	2.388535	12.56
2.	2.260739	13.27	2.259036	13.28	2.320186	12.93	2.281369	13.15
3.	2.431118	12.34	2.331002	12.87	2.302379	13.03	2.288330	13.11
Average	2.339747	12.83	2.298652	13.05	2.335136	12.85	2.319411	12.94
Standard Deviation	0.070104	0.38	0.029826	0.17	0.034508	0.188	0.048960	0.270

The results obtained of GIOPS and GFLOPS are visualized in the below chart:

CPU (Average)



CPU(Std. Deviation)



Efficiency:

Efficiency can be calculated as percentage of actual performance by theoretical performance.

$$\text{Efficiency} = \frac{\text{Actual Performance}}{\text{Theoretical Performance}} * 100$$

Thus, efficiency in GFLOPS can be calculated as,


$$= (2.3397/36.64) * 100$$

$$= \mathbf{6.39\%}$$


CPU BENCHMARKING FOR 600 SAMPLES:

- As a separate experiment, the benchmark code is run on double precision floating point and integer instructions, 8 threads for a 10-minute period for each one, and samples were taken every second on how many instructions per second were achieved during the experiment.
- There are separate code files written for integer instructions and double precision floating point instructions.
- The 600 samples are recorded and the values are plotted with time (0 to 10 min) on the x-axis and FLOPS/IOPS on the y-axis.

CPU (600 samples)



CPU (600 samples)



LINPACK BENCHMARK:

- The Linpack benchmark was run across all the nodes of the instance VM and the best GFLOPS performance was recorded.
- The efficiency in the speed achieved compared to both theoretical performance calculated and the benchmark results obtained from our code is reported.

The screenshot below shows the results of Linpack Benchmark run on Chameleon openstack instance:

```

root@pa1-gudur-reddy-bk5-inst:/home/cc/pa1/CPU
Input data or print help ? Type [data]/help :
data
Number of equations to solve (problem size): 10000
Leading dimension of array: 10000
Number of trials to run: 4
Data alignment value (in Kbytes): 4
Current date/time: Sun Oct 8 02:53:45 2017

CPU frequency: 1.479 GHz
Number of CPUs: 2
Number of cores: 2
Number of threads: 1

Parameters are set to:

Number of tests: 1
Number of equations to solve (problem size) : 10000
Leading dimension of array : 10000
Number of trials to run : 4
Data alignment value (in Kbytes) : 4

Maximum memory requested that can be used=800204096, at the size=10000

===== Timing linear equation system solver =====
Size LDA Align. Time(s) GFlops Residual Residual(norm) Check
10000 10000 4 18.710 35.6415 6.859494e-11 2.418727e-02 pass
10000 10000 4 18.575 35.9021 6.859494e-11 2.418727e-02 pass
10000 10000 4 19.213 34.7088 6.859494e-11 2.418727e-02 pass
10000 10000 4 18.458 36.1286 6.859494e-11 2.418727e-02 pass

Performance Summary (GFlops)
Size LDA Align. Average Maximal
10000 10000 4 35.5953 36.1286

Residual checks PASSED

End of tests

[root@pa1-gudur-reddy-bk5-inst CPU]#

```

The parameters set are as follows:

Number of CPUs : 2
Number of Cores : 2
Number of threads : 1

The average peak performance for a problem size of 10000 was found to be **35.5953 GFLOPS**.

Efficiency achieved = $(\text{Linpack performance} / \text{Theoretical speed}) * 100$

$$= (35.5953 / 36.64) * 100$$

$$= \mathbf{97.15 \%}$$

Efficiency achieved by Linpack benchmark is much higher than efficiency obtained by our code (6.39%) and is close to theoretical performance obtained.

Thus, we conclude that Linpack benchmarks are much better than those obtained in our benchmark. As per Netlib.org, number of factors affect the performance of a processor including problem size, load on the system, accuracy of the clock, compiler options, version of the

compiler, size of cache, bandwidth from memory, amount of memory, etc. The problem size used in Linpack is 10000 which is 10^5 times lesser than in our benchmark. Also, Linpack benchmark code uses matrix multiplications which increases the performance by a substantial amount. Linpack itself is built on BLAS, which is a high performance package for computations. Linpack also makes use of AVX instructions which boosts the speed.

3. Conclusion :

Thus, speed increase with increase in number of threads.

The performance achieved by our code was below theoretical speed obtained. Number of factors are responsible : lack of use of AVX instructions being one of them.

Linpack achieved the best CPU speed which is almost equal to theoretical speed.

4. Improvements :

1. GUI can be implemented for real-time tracking of processor speed.
2. Complex operations involving matrix multiplications

MEMORY BENCHMARKING:

1. Program Design

The Memory benchmarking program is designed as follows.

The program takes a user input of the block size and the number of threads to run the program on.

It consists of 16 functions, out of which, 4 functions correspond to the block size and is used for the calculation of the latency and throughput.

The other functions (4*3) = 12 functions for read+write and sequential write and random write (using memset(), memcpy()).

The last function is createFile() which is used to create a 1GB file to perform the operations on.

2. Performance

The throughput and latency values obtained are as follows:

Throughput - 1 thread (MBps):

SNo	Read + Write				Sequential Write				Random Write			
	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB
1	800	3289.473	7876.923	2782.608	800	3906.250	20480	4000	114.28	2717.391	4654.545	4000
2	800	3289.473	7314.285	2909.090	400	3906.250	17066.67	4000	133.33	2840.909	4876.1904	4571.428
Avg	800	3289.473	5393	2845.85	600	3906.250	18773.34	4000	123.80	2779.15	4765.377	4285.784
SD	0	0	281.319	63.241	200	0	1706.665	0	9.525	61.759	110.8227	285.714

Latency - 1 thread (usec) :

SNo	Read + Write				Sequential Write				Random Write			
	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB
1	1250	304	126.95 3	359.375	1250	256	48.83	250	0.0000009	368	214.84 3	250
2	1250	304	136.71 8	343.750	2500	256	58.60	250	0.0000007	352	205.07 8	218.750
Avg	1250	304	131.84	351.562	1875	256	53.72	250	0.0000000 8	360	209.96	234.375
SD	0	0	4.8825	7.8125	625	0	4.885	0	0	8	4.8825	15.625

Throughput - 2 threads (MBps):

SNo	Read + Write				Sequential Write				Random Write			
	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB
1	800	3289. 473	6826.6 7	2909.09 0	400	3676. 470	17066 .67	4266.67	133.33	2040. 909	4654.5 45	3764.705
2	800	3289. 473	6826.6 7	3047.61 9	800	3676. 470	20480	4000	114.285	2976. 190	4654.5 45	4266.67
Avg	800	3289. 473	6826.6 7	2978.35 5	600	3676. 470	18773 .34	4133.34	123.80	2508. 55	4654.5 45	4015.69
SD	0	0	0	69.2645	200	0	1706. 665	133.335	9.5225	467.6 405	0	250.9825

Latency - 2 threads (usec) :

SNo	Read + Write				Sequential Write				Random Write			
	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB

1	1250	304	146	343.75	2500	272	58.593	234.375	0.00000007	352	214.843	265.625
2	1250	304	146	328.125	1250	272	48.82	250	0.00000009	336	214.843	234.375
Avg	1250	304	146	336	1875	272	53.70	242.19	0.000000085	344	214.843	250
SD	0	0	0	7.8125	625	0	4.885	7.8125	0	8	0	15.625

Throughput - 4 threads (MBps):

SNo	Read + Write				Sequential Write				Random Write			
	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB
1	800	3289.473	7314.285	2909.090	800	3289.473	20480	4000	100	2604.16	4452.174	4000
2	800	3125	7314.285	2782.608	400	3906.25	17066.67	4000	133.33	2976.190	4452.174	3764.705
Avg	800	3207.23	7314.285	2845.849	600	3597.861	18773.335	4000	116.67	2790.175	4452.174	3882.3585
SD	0	82.2365	0	63.241	200	308.3885	1706.665	0	16.665	186.015	0	117.6475

Latency - 4 threads (usec) :

SNo	Read + Write				Sequential Write				Random Write			
	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB
1	1250	304	136.718	343.75	1250	304	48.828	250	0.000000010	384	224.609	250
2	1250	320	136.718	359.375	2500	256	585.93	250	0.00000007	336	224.609	265.62
Avg	1250	312	136.718	351.56	1875	280	317.38	250	0.000000085	360	224.609	257.81
SD	0	8	0	7.8125	625	24	268.551	0	0	24	0	7.81


Throughput - 8 threads (MBps):

SNo	Read + Write				Sequential Write				Random Write			
	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB
1	800	3125	7314.285	2909.090	400	3906.250	17066.67	4000	114.285	2840.909	4654.545	4000
2	800	3125	6023.529	2782.608	400	3906.250	14628.571	4000	133.333	3289.273	5120	3764.705
Avg	800	3125	6668.907	2845.85	400	3906.250	15847.62	4000	123.809	3065.091	4887.272	3882.352
SD	0	0	645.378	63.241	0	0	1219.0495	0	9.524	224.182	232.7275	117.647


Latency - 8 threads (usec) :

SNo	Read + Write				Sequential Write				Random Write			
	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB
1	1250	320	136.718	343.750	2500	256	58.593	250	0.00000009	352	214.843	250
2	1250	320	166.015	359.375	2500	256	68.36	250	0.00000007	304	195.312	265.625
Avg	1250	320	151.367	351.562	2500	256	63.48	250	0.000000085	328	205.078	257.813
SD	0	0	14.6485	7.8125	0	0	4.8835	0	0	24	9.7655	7.8125


Read + Write (Throughput)



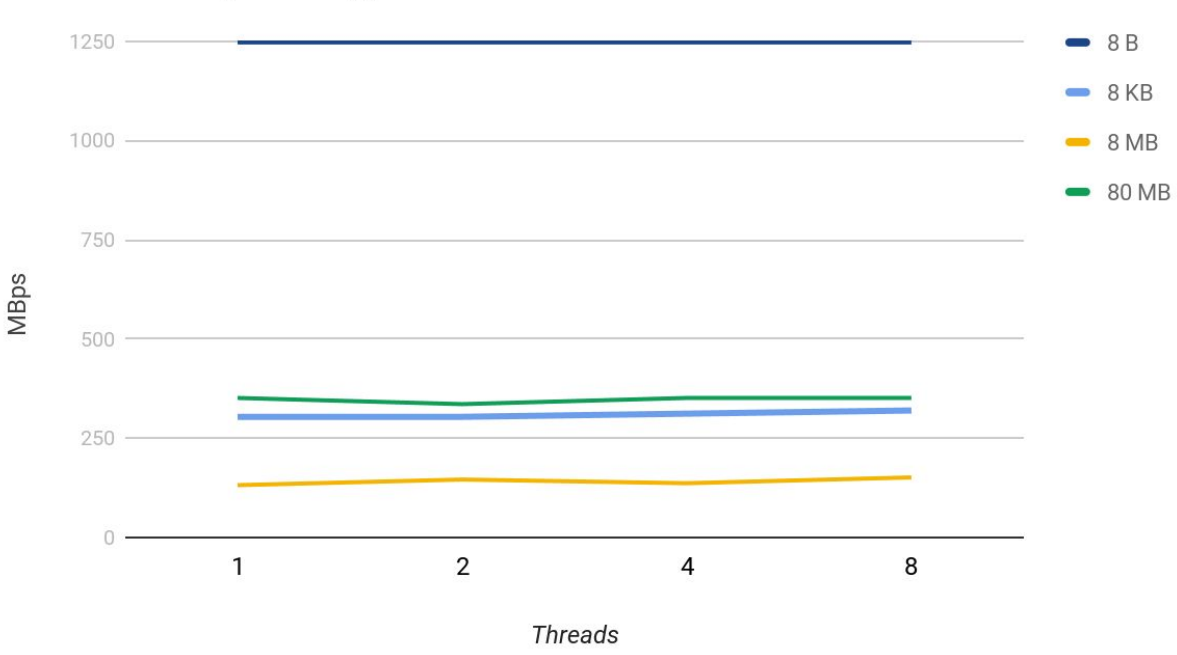
Sequential Write (Throughput)




Random Write (Throughput)



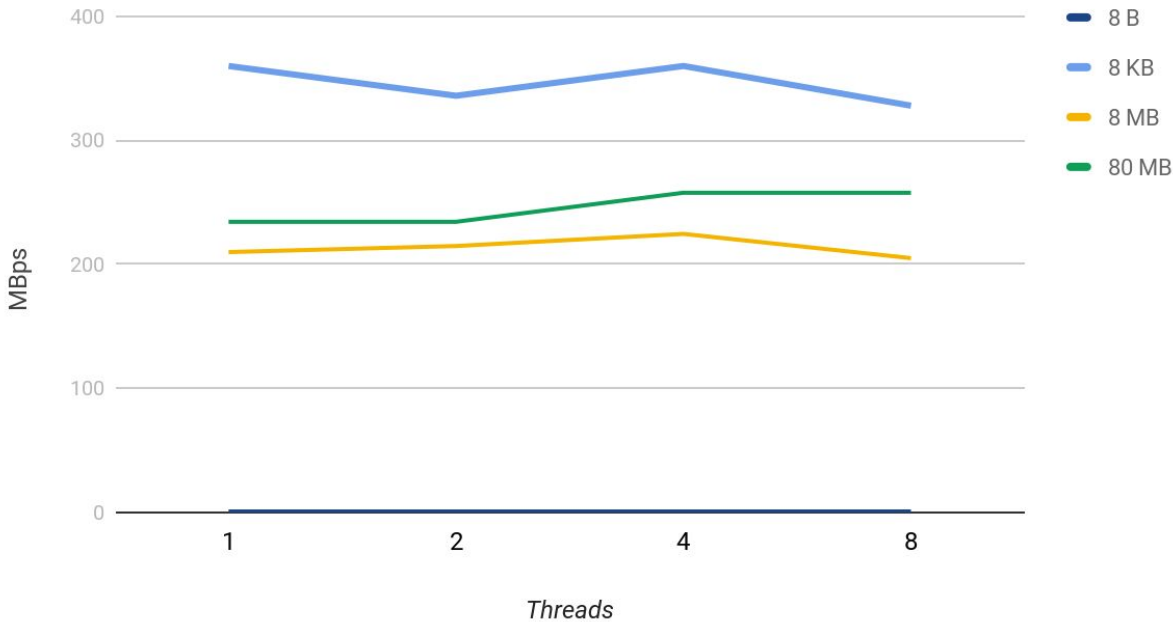
Read + Write (Latency)



Sequential Write (Latency)



Random Write (Latency)



Theoretical Memory Bandwidth :

Theoretical memory bandwidth can be roughly calculated as ,

$$\begin{aligned} &= (2 * \text{dimmSpeed of RAM(DDR3)} * \text{max_width of the packet})/8 \\ &= (2 * 800\text{Mhz} * 64)/8 \\ &= \mathbf{12800 \text{ MBps}} \end{aligned}$$

From the Hardware description of the Chameleon cloud, the network bandwidth is known as 100Gbps=12500MBps

$$\begin{aligned} \text{Efficiency achieved} &= (7314.285/12500) * 100 \\ &= \mathbf{58.514\%} \end{aligned}$$

STREAM BENCHMARK:

The stream benchmark was run on openkvm instance. The screenshot of the results are as follows:

```
[cc@pal-gudur-reddy MEMORY]$ gcc -O stream.c -o stream
[cc@pal-gudur-reddy MEMORY]$ ./stream
-----
STREAM version $Revision: 5.10 $
-----
This system uses 8 bytes per array element.
-----
Array size = 10000000 (elements), Offset = 0 (elements)
Memory per array = 76.3 MiB (= 0.1 GiB).
Total memory required = 228.9 MiB (= 0.2 GiB).
Each kernel will be executed 10 times.
The *best* time for each kernel (excluding the first iteration)
will be used to compute the reported bandwidth.
-----
Your clock granularity/precision appears to be 1 microseconds.
Each test below will take on the order of 12228 microseconds.
(= 12228 clock ticks)
Increase the size of the arrays if this shows that
you are not getting at least 20 clock ticks per test.
-----
WARNING -- The above is only a rough guideline.
For best results, please be sure you know the
precision of your system timer.
-----
Function      Best Rate MB/s  Avg time     Min time     Max time
Copy:         11474.5   0.014037     0.013944     0.014410
Scale:        11464.5   0.014109     0.013956     0.014467
Add:          12730.7   0.019025     0.018852     0.019244
Triad:        12412.1   0.019898     0.019336     0.020312
-----
Solution Validates: avg error less than 1.000000e-13 on all three arrays
[cc@pal-gudur-reddy MEMORY]$
```

Thus, the best performance obtained using stream benchmark is **11474 MB/s** using stream benchmark.

$$\begin{aligned} \text{Efficiency} &= (\text{stream performance/theoretical throughput}) * 100 \\ &= (11474/12500) * 100 \\ &= \mathbf{91.79\%} \end{aligned}$$

Thus, the efficiency obtained by stream benchmark is high, close to theoretical value.

3. Conclusion :

Thus, best performance is achieved at concurrency level 4.

Also, the sequential access is faster than the random access because the overhead of accessing any random memory and seeking the file pointer to that random location is avoided.

4. Improvements :

1. GUI can be implemented.
2. Efficient memory handling functions can be implemented.

DISK BENCHMARKING:

1. Program Design

The Disk program is designed as follows.

The program takes a user input of the block size and the number of threads to run the program on.

It consists of 17 functions, out of which, 4 functions correspond to the block size and is used for the calculation of the latency and throughput.

The other functions (4*3) = 12 functions for read+write and sequential read and random read(using fread(), fwrite() and fseek()).

The last function is createFile() which is used to create a 10GB file to perform the operations on.

2. Performance

The throughput and latency values obtained are as follows:

Throughput - 1 thread (MBps):

SNo	Read + Write				Sequential Read				Random Read			
	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB
1	9.7656 25	64.87 3	410.25 64	49.2610	19.53 125	98.39	1777.7 78	49.81 32	26.04 1667	86.38	941.1 764	49.8132
2	11.38	59.3 93	489.9 782	48.134	22.94	93.0 6	1834.9	47.84 6	30.72 8	85.57	1067. 298	46.948
Avg	10.5 7	62.1 33	450.1 141	48.6975	21.24	95.7 25	1806. 339	54.15	28.38	85.975	1004. 237	48.3806

Throughput - 2 threads (MBps):

SNo	Read + Write				Sequential Read				Random Read			
	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB
1	11.95	78.125	470.588	46.7836	21.478	83.6893	761.904	48.721	35.894	78.125	1066.667	49.93757
2	11.38	74.987	493.2963	48.305	19.383	80.398	794.37	52.7856	29.399	74.834	1205.76823	54.89342
Avg	11.67	76.556	481.942	47.5443	20.43	82.04365	778.137	50.7533	32.64	76.4795	1136.217	52.41546

Throughput - 4 threads (MBps):


SNo	Read + Write				Sequential Read				Random Read			
	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB
1	8.373	78.125	457.14	46.457	16.389	83.6893	800.00	46.403	10.38	78.125	1333.3	49.200
2	5.278	74.987	386.4528	43.7289	18.389	80.398	639.924	43.562	7.378	74.834	1194	49.458
Avg	6.83	76.556	421.7964	45.092	17.389	82.04365	719.962	44.9825	8.88	76.4795	1263.665	49.329

Throughput - 8 threads (MBps) :


SNo	Read + Write				Sequential Read				Random Read			
	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB	8 B	8 KB	8 MB	80 MB
1	10.390	78.125	457.14	47.675	16	49.543	800.00	49.87	9	39.0625	1333.3	49.937

2	10	72.94 4	429.43 9	50.435	15.87	42.95 34	753.8 92	52.89	9.34	36.767 4	1274.4 510	51.34
Avg	10.2	75.53 45	443.28 95	49.055	15.43 5	46.24 82	776.9 5	51.38	9.17	37.914 7	1303.8 755	50.6385

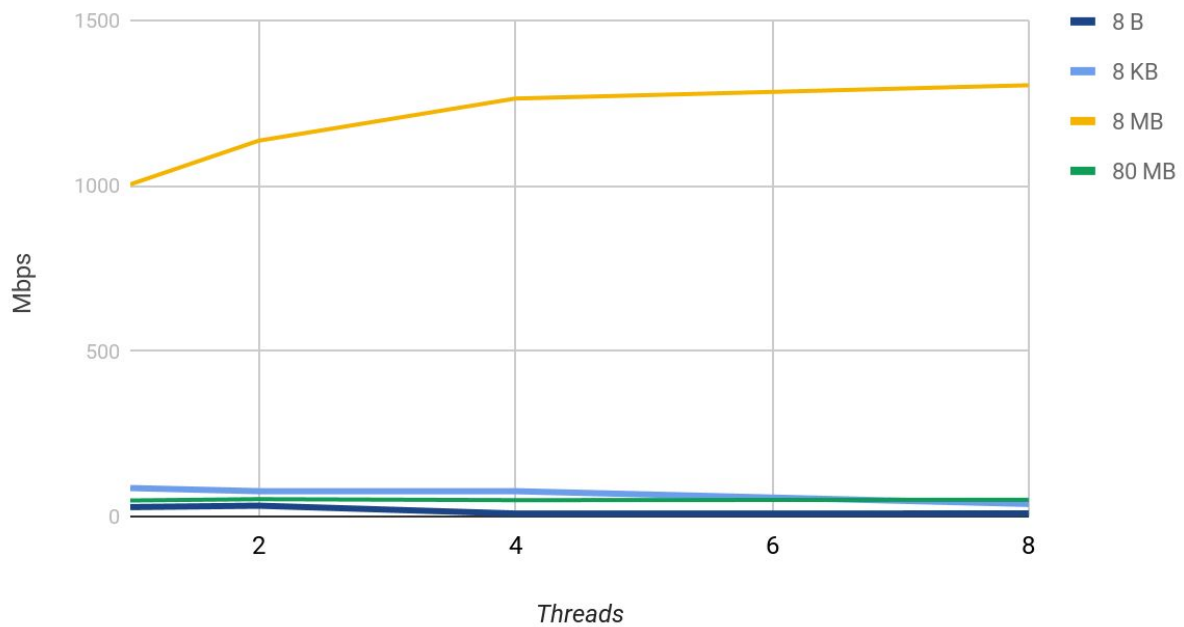
Read + Write (Throughput)



Sequential Read (Throughput)



Random Read (Throughput)



Thus, Random Read and Sequential Read values are almost similar in a Solid State Drive (SSD). Since the values of Random Read and Sequential Read obtained above are

similar, we can say that the disk being evaluated is an SSD. The optimal performance was obtained at the concurrency level 2.

Theoretical performance given :

= 6TB inline SAS Drive

Throughput : 6Gbps (Read + Write)

$$\begin{aligned}\text{Efficiency achieved} &= \frac{(\text{Actual Performance})}{(\text{Theoretical Performance})} * 100 \\ &= (3.76/6) * 100 \quad (\text{Since } 470\text{MBps} = 3.76\text{Gbps}) \\ &= \mathbf{62.67\%}\end{aligned}$$

IOZONE BENCHMARK:

The result obtained after performing iotzone benchmark for a file size is 8 MB is **271077 KB/s which is 2.168616 Gbps.**

```
[cc@pa1-gudur-reddy-bk6-1 current]$ ./iozone -a -r 8M -s 10G -T
Iotzone: Performance Test of File I/O
Version $Revision: 3.471 $
Compiled for 64 bit mode.
Build: linux

Contributors:William Norcott, Don Capps, Isom Crawford, Kirby Collins
Al Slater, Scott Rhine, Mike Wisner, Ken Goss
Steve Landherr, Brad Smith, Mark Kelly, Dr. Alain CYR,
Randy Dunlap, Mark Montague, Dan Million, Gavin Brebner,
Jean-Marc Zucconi, Jeff Blomberg, Benny Halevy, Dave Boone,
Erik Habbinga, Kris Strecker, Walter Wong, Joshua Root,
Fabrice Bacchella, Zhenghua Xue, Qin Li, Darren Sawyer,
Vangel Bojaxhi, Ben England, Vikentsi Lapa,
Alexey Skidanov.

Run began: Mon Oct 9 16:10:24 2017

Auto Mode
Record Size 8192 kB
File size set to 10485760 kB
Command line used: ./iozone -a -r 8M -s 10G -T
Output is in kBytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 kBytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.

      kB  reclen  write  rewrite  read  reread  random  random  bkwd  record  stride
10485760  8192  139435  271077  240479  272543  217201  160464  140163  8290533  141783  213304  135893  164936  175763

iozone test complete.
```

$$\begin{aligned}\text{Efficiency achieved} &= (2.168/6) * 100 \\ &= \mathbf{36.13 \%}\end{aligned}$$

3. Conclusion

- The throughput is the highest for the 8MB block size.
- The Read + write is slower than random and sequential Read. Also, the Sequential read is higher than the speed of random read. This is because of the fact that the disk head should move more for the random case in a hard disk. Hence the difference in the speeds of random and sequential reads is less in the case of a SSD.

4. Improvements and Extensions

- While the test is in progress, for optimal results, the system should be held idle.
- If the code is written more efficiently by using a better API, better results can be achieved.
- A GUI increases the ease of use of the benchmarking test.

NETWORK BENCHMARKING:

1. Program Design

The main goal of this experiment is to establish TCP/UDP connection between client and server and measure the network speed between one node as well as two nodes.

The metrics used for measurement are throughput and latency, at varying levels of currency.

- The design structure of the program is such that separate codes are written for both TCP and UDP.
- Both client and server send and receive data of the size of 64KB, 100 times.
- The throughput is measured the number of loops times data size divided by the time required to send and receive the data, in Mbits per sec.
- Latency is measured as execution time divided by number of iterations.
- Strong scaling is implemented at both server side and client side.

2. Performance

This experiment was performed on openkvm instance. The results obtained after performing the experiment several times are as follows:

TCP:

1 Node:


Sr. No.	1 thread		2 threads		4 threads		8 threads	
	Throughput	latency	Throughput	latency	Throughput	latency	Throughput	latency
1.	3759.681	0.133	3681.342	0.136	3527.336	0.142	3377.693	0.148
2.	3683.783	0.136	3963.221	0.126	4082.299	0.122	3910.833	0.128
3.	3885.607	0.129	3691.399	0.135	3603.343	0.138	3843.392	0.130
Average	3776.357	0.132	3778.654	0.132	3737.659	0.134	3710.639	0.135

Standard Deviation	83.233	0.0028	130.573	0.0045	245.664	0.0086	237.033	0.0089
---------------------------	---------------	---------------	----------------	---------------	----------------	---------------	----------------	---------------


2 Nodes:

Sr. No.	1 thread		2 threads		4 threads		8 threads	
	Throughput	latency	Throughput	latency	Throughput	latency	Throughput	latency
1.	550.964	0.908	498.529	1.003	528.306	0.946	609.377	0.820
2.	481.245	1.039	515.533	0.969	536.032	0.932	530.571	0.942
3.	500.265	1.000	503.829	0.992	562.809	0.888	532.305	0.929
Average	510.824	0.982	505.96	0.988	542.382	0.922	557.417	0.897
Standard Deviation	29.42	0.0549	7.104	0.0141	14.784	0.0247	36.747	0.0547

TCP (Average)



TCP (Std deviation)



UDP:


1 Node:

Sr. No.	1 thread		2 threads		4 threads		8 threads	
	Throughput	latency	Throughput	latency	Throughput	latency	Throughput	latency
1.	9403.798	0.0531	9516.588	0.0525	8484.642	0.0589	8271.898	0.0604
2.	9433.961	0.0530	9602.358	0.0520	9380.863	0.0533	8264.462	0.0605
3.	10.864.840	0.0460	9237.021	0.0541	9031.792	0.0553	8298.754	0.0602
Average	9418.8795	0.0507	9451.989	0.0528	8965.765	0.0558	8278.371	0.0603
Standard Deviation	15.081	0.0033	155.986	0.00089	368.847	0.00231	14.728	0.000124


2 Nodes:

Sr. No.	1 thread		2 threads		4 threads		8 threads	
	Throughput	latency	Throughput	latency	Throughput	latency	Throughput	latency
1.	9442.870	0.0529	9267.840	0.0539	8716.875	0.0573	8385.0410	0.0596
2.	9543.805	0.0523	6434.178	0.0771	2668.089	0.1874	8272.667	0.0604
3.	9209.799	0.0543	9281.603	0.0538	9196.249	0.0543	8638.562	0.0578
Average	9398.824	0.05316	8327.873	0.0616	6860.404333	0.0996	8432.09	0.0592
Standard Deviation	139.868	0.000837	1339.056	0.01096	2970.86	0.06204	153.035	0.00108

UDP (Average)



UDP (Std Deviation)



For TCP, Thus, the best throughput performance, 1 node obtained over the loopback interface is **3776.357 Mbits/s , 3778.654 MBits/s, 3737.659 MBits/s, 3710.639 MBits/s** for **1, 2 , 4 and 8 threads** respectively.

For 2 nodes, throughput performance obtained is **510.824 Mbits/s , 505.96 MBits/s, 542.382 MBits/s, 557.417 MBits/s** for **1,2,4, and 8 threads** respectively.

For UDP, Thus, the best throughput performance, 1 node obtained over the loopback interface is **9418.8795 MBits/s, 9451.989 Mbits/s, 8965.765 MBits/s and 8278.371 MBits/s** for **1, 2 , 4 and 8 threads** respectively.

For 2 nodes, throughput performance obtained is **9398.824 MBits/s, 8327.873 MBits/s, 6860.404333 MBits/s, 8432.09 MBits/s** for **1,2,4, and 8 threads** respectively.

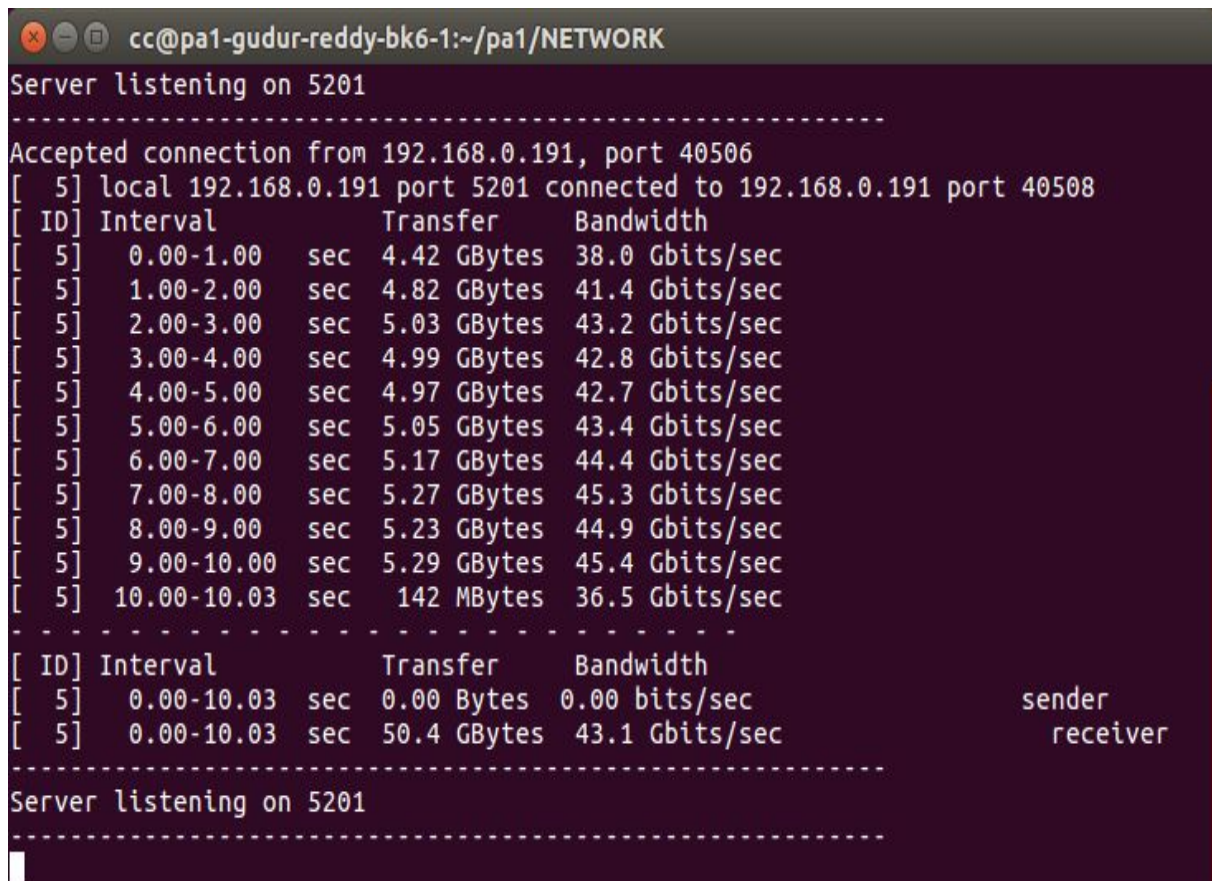
Theoretical Network Speed = 12500 MBits/s

IPerf Benchmark:

IPerf benchmark is run on two separate Baremetal instances and performance is recorded.

The screenshots of client and server are given below:

1 Node:



```
cc@pa1-gudur-reddy-bk6-1:~/pa1/NETWORK
Server listening on 5201
-----
Accepted connection from 192.168.0.191, port 40506
[ 5] local 192.168.0.191 port 5201 connected to 192.168.0.191 port 40508
[ ID] Interval           Transfer     Bandwidth
[ 5]  0.00-1.00   sec   4.42 GBytes  38.0 Gbits/sec
[ 5]  1.00-2.00   sec   4.82 GBytes  41.4 Gbits/sec
[ 5]  2.00-3.00   sec   5.03 GBytes  43.2 Gbits/sec
[ 5]  3.00-4.00   sec   4.99 GBytes  42.8 Gbits/sec
[ 5]  4.00-5.00   sec   4.97 GBytes  42.7 Gbits/sec
[ 5]  5.00-6.00   sec   5.05 GBytes  43.4 Gbits/sec
[ 5]  6.00-7.00   sec   5.17 GBytes  44.4 Gbits/sec
[ 5]  7.00-8.00   sec   5.27 GBytes  45.3 Gbits/sec
[ 5]  8.00-9.00   sec   5.23 GBytes  44.9 Gbits/sec
[ 5]  9.00-10.00  sec   5.29 GBytes  45.4 Gbits/sec
[ 5] 10.00-10.03  sec    142 MBytes  36.5 Gbits/sec
-----
[ ID] Interval           Transfer     Bandwidth
[ 5]  0.00-10.03  sec    0.00 Bytes   0.00 bits/sec
[ 5]  0.00-10.03  sec   50.4 GBytes  43.1 Gbits/sec
-----
Server listening on 5201
-----
```

sender
receiver

Fig : server side

```

vaishnavi@vaishnavi-VirtualBox:~$ ssh -i cloud.key cc@129.114.33.25
Last login: Mon Oct  9 18:14:48 2017 from 68.61.156.76
[cc@pa1-gudur-reddy-bk6-1 ~]$ cd pa1/NETWORK/
[cc@pa1-gudur-reddy-bk6-1 NETWORK]$ sudo rpm -Uvh iperf3-3.1.3-1.fc24.x86_64.rpm

warning: waiting for transaction lock on /var/lib/rpm/.rpm.lock
Preparing... ##### [100%]
package iperf3-3.1.3-1.fc24.x86_64 is already installed
[cc@pa1-gudur-reddy-bk6-1 NETWORK]$ clear
[cc@pa1-gudur-reddy-bk6-1 NETWORK]$ iperf3 -c 192.168.0.191
Connecting to host 192.168.0.191, port 5201
[ 4] local 192.168.0.191 port 40508 connected to 192.168.0.191 port 5201
[ ID] Interval          Transfer      Bandwidth    Retr  Cwnd
[ 4]  0.00-1.00    sec   4.54 GBytes   39.0 Gbits/sec    0   2.12 MBytes
[ 4]  1.00-2.00    sec   4.83 GBytes   41.5 Gbits/sec    0   2.12 MBytes
[ 4]  2.00-3.00    sec   5.03 GBytes   43.2 Gbits/sec    0   2.12 MBytes
[ 4]  3.00-4.00    sec   4.99 GBytes   42.8 Gbits/sec    0   2.12 MBytes
[ 4]  4.00-5.00    sec   4.96 GBytes   42.6 Gbits/sec    0   2.12 MBytes
[ 4]  5.00-6.00    sec   5.06 GBytes   43.5 Gbits/sec    0   2.12 MBytes
[ 4]  6.00-7.00    sec   5.16 GBytes   44.3 Gbits/sec    0   2.12 MBytes
[ 4]  7.00-8.00    sec   5.29 GBytes   45.4 Gbits/sec    0   2.12 MBytes
[ 4]  8.00-9.00    sec   5.22 GBytes   44.9 Gbits/sec    0   2.12 MBytes
[ 4]  9.00-10.00   sec   5.29 GBytes   45.5 Gbits/sec    0   2.12 MBytes
- - - - -
[ ID] Interval          Transfer      Bandwidth    Retr
[ 4]  0.00-10.00   sec   50.4 GBytes   43.3 Gbits/sec    0
[ 4]  0.00-10.00   sec   50.4 GBytes   43.3 Gbits/sec    0
sender
receiver

iperf Done.
[cc@pa1-gudur-reddy-bk6-1 NETWORK]$ █

```

Fig : client side

$$\begin{aligned}
 \text{Efficiency achieved} &= (4330/12500) * 100 \\
 &= \mathbf{34.64\%}
 \end{aligned}$$

2 Nodes:

```
cc@pa1-gudur-reddy-bk6:~/pa1/NETWORK
Server listening on 5201
-----
Accepted connection from 192.168.0.191, port 37410
[ 5] local 192.168.0.190 port 5201 connected to 192.168.0.191 port 37412
[ ID] Interval      Transfer    Bandwidth
[ 5]  0.00-1.00    sec      194 MBytes  1.63 Gbits/sec
[ 5]  1.00-2.00    sec      300 MBytes  2.52 Gbits/sec
[ 5]  2.00-3.00    sec      304 MBytes  2.55 Gbits/sec
[ 5]  3.00-4.00    sec      310 MBytes  2.60 Gbits/sec
[ 5]  4.00-5.00    sec      295 MBytes  2.47 Gbits/sec
[ 5]  5.00-6.00    sec      299 MBytes  2.51 Gbits/sec
[ 5]  6.00-7.00    sec      305 MBytes  2.56 Gbits/sec
[ 5]  7.00-8.00    sec      303 MBytes  2.54 Gbits/sec
[ 5]  8.00-9.00    sec      293 MBytes  2.46 Gbits/sec
[ 5]  9.00-10.00   sec      301 MBytes  2.53 Gbits/sec
[ 5] 10.00-10.04   sec      11.1 MBytes 2.47 Gbits/sec
-----
[ ID] Interval      Transfer    Bandwidth
[ 5]  0.00-10.04   sec      0.00 Bytes  0.00 bits/sec
[ 5]  0.00-10.04   sec      2.85 GBytes 2.44 Gbits/sec
-----
Server listening on 5201
-----
```

Fig : Server side instance

```
cc@pa1-gudur-reddy-bk6-1:~/pa1/NETWORK
Report bugs to: https://github.com/esnet/iperf
[cc@pa1-gudur-reddy-bk6-1 NETWORK]$ clear

[cc@pa1-gudur-reddy-bk6-1 NETWORK]$ iperf3 -c 192.168.0.190
Connecting to host 192.168.0.190, port 5201
[ 4] local 192.168.0.191 port 37412 connected to 192.168.0.190 port 5201
[ ID] Interval      Transfer    Bandwidth    Retr    Cwnd
[ 4]  0.00-1.00    sec      208 MBytes  1.75 Gbits/sec  19    387 KBytes
[ 4]  1.00-2.00    sec      301 MBytes  2.52 Gbits/sec   3    539 KBytes
[ 4]  2.00-3.00    sec      304 MBytes  2.55 Gbits/sec   9    522 KBytes
[ 4]  3.00-4.00    sec      310 MBytes  2.60 Gbits/sec   8    525 KBytes
[ 4]  4.00-5.00    sec      294 MBytes  2.47 Gbits/sec   7    505 KBytes
[ 4]  5.00-6.00    sec      299 MBytes  2.51 Gbits/sec  18    482 KBytes
[ 4]  6.00-7.00    sec      306 MBytes  2.57 Gbits/sec  13    486 KBytes
[ 4]  7.00-8.00    sec      303 MBytes  2.54 Gbits/sec   9    484 KBytes
[ 4]  8.00-9.00    sec      293 MBytes  2.46 Gbits/sec   7    462 KBytes
[ 4]  9.00-10.00   sec      301 MBytes  2.53 Gbits/sec  16    464 KBytes
-----
[ ID] Interval      Transfer    Bandwidth    Retr
[ 4]  0.00-10.00   sec      2.85 GBytes 2.45 Gbits/sec  109
[ 4]  0.00-10.00   sec      2.85 GBytes 2.45 Gbits/sec
-----
iperf Done.
[cc@pa1-gudur-reddy-bk6-1 NETWORK]$
```

Fig : Client side instance

Thus, the best performance achieved by IPerf benchmark is **2.45 Gbits/sec** i.e 2450 Mbits/sec.

$$\text{Efficiency achieved} = \frac{(\text{Actual performance})}{(\text{Theoretical Performance})} * 100$$

$$\begin{aligned} &= (2450/12500) * 100 \\ &= \mathbf{20 \%} \end{aligned}$$

3. Conclusion :

Thus, network speed achieved is greater on a loopback interface on one node as compared to that achieved on two nodes. This may be due to a variety of factors including network traffic, congestion.

UDP transmission gives greater network speed however, it is unreliable compared to TCP.

IPerf benchmark achieved greater efficiency on one node.

4. Improvements and Extensions:

1. GUI can be implemented.
2. Varying size of data can be transmitted and the program can be tested on all levels to observe the deviation.